



User Guide for Delay/Disruption Tolerant Networking Bundle Protocol Node (BPNode) v7.0.5

NASA GSFC DTN Project 450.2-DTN-UG

Version 2.0, March 2026

Table of Contents

Signatures	1
Change History	2
1. Introduction	3
1.1. Purpose and Scope	3
1.2. Document Review, Approval, and Update	3
1.3. Support	3
1.4. Assumptions	4
2. Software Overview	5
2.1. Supported Platforms	5
2.2. cFS Framework	5
2.3. BPNode and BPLib	7
2.4. Dependencies	8
2.5. QCBOR	9
3. Installation	10
3.1. Building BPNode and BPLib within cFS	10
3.1.1. Clone Repositories	10
3.1.2. Adding BPNode Configurations to cFS	10
3.1.2.1. dtn_defs Content	11
3.1.3. Building and Deploying the DTN cFS	12
3.1.4. Run Unit Tests and Generate Coverage	12
3.1.5. Troubleshooting	13
3.2. Standalone BPLib Build (bpcat)	13
4. Configuration	15
4.1. Startup Configuration	15
4.1.1. ADU Channels	15
4.1.2. CLA Contacts	16
5. Running the Node	17
5.1. Managing Channel Applications	18
5.2. Managing Contacts	19
5.3. Bundle Storage Management	20
5.4. Running a Sample Bundle Flow	20
6. Directives, Telemetry, Tables, and Events	23
7. Custody Transfer	24
7.1. Bundle Definitions	24
7.2. Data Flow	25
7.2.1. When to Accept or Reject Custody	26
7.2.2. Custodial Bundle Handling	27
7.2.3. Custody Signal Processing Flow	28

7.3. Configuration	29
7.4. Operation	29
8. Known Issues	31
9. DTN Tools Suite	32
9.1. DTNGEN	32
9.2. DTNCLA	35
9.3. Dependencies	35
9.4. Tools Suite Installation	35
9.5. COSMOS Integration	36
Appendix A: COSMOS Interface	38
Appendix B: Host Assumptions	40
Appendix C: Acronyms	43

Signatures

Prepared by:

Sara Garcia-Beech
DTN FSW Product Development Lead
Flight Software Systems Branch
Code 534, NASA GSFC

Date

Rebecca Besser
DTN Systems Engineering Lead and Acting Project Manager
Software Systems Engineering Branch
Code 531, NASA GSFC

Date

Change History

Version	Date	Description	Affected Pages
1.0	09/2025	Initial Release	All
2.0	03/2026	Update for 7.0.5 release	All

Chapter 1. Introduction

The Bundle Protocol (BP) application node (BPNode) is a core Flight System (cFS) application that provides RFC-9171 BP services to embedded space flight systems by integrating BP library (BPLib) with cFS framework and processes. The Goddard Space Flight Center (GSFC) Delay/Disruption Tolerant Networking (DTN) Project designed and developed BPNode as a generic flight-capable node compliant with Bundle Protocol version 7 (BPv7).

BPNode was developed in accordance with NASA Systems Engineering Processes and Requirements (NPR 7123.1D) and Software Class B requirements as defined in NASA Software Engineering Requirements (NPR 7150.2D). It relies on a subset of cFS utilities (such as fault handling, task scheduling, data recording, etc.) to avoid duplication of code and leverage a software framework common to many human and robotic missions. This framework allows BPNode to run without modifications on various hardware and software configurations. Missions can tailor their operational environment based on available cFS configurations.

BPNode also complies with the Consultative Committee for Space Data Systems (CCSDS) standards for international interoperability and conforms with the LunaNet Interoperability Specification. BPNode enables cFS to communicate with other DTN nodes using Bundle Protocol.

1.1. Purpose and Scope

This BPNode User Guide provides its users with a general understanding of node functions. It covers node installation, configuration, and operations. The guide is intentionally generic to enable customization.

This software is available as an open source distribution, released under the Apache 2.0 license. The open-source license enables the industry adoption of BPNode. Outside of major releases, updates to the open repositories are limited to bug fixes and minor enhancements to these components.

1.2. Document Review, Approval, and Update

This guide is developed and maintained by the DTN Project at NASA Goddard Space Flight Center. The document has been reviewed and approved by the team members as listed on the signature page. After initial approval, the document will be treated as a controlled document, placed under configuration management. Changes are listed in the document [Change History](#) log.

1.3. Support

The development team also offers training, support, and subject matter expertise to government and industry users, and can provide custom development and integration services for flight projects. A support agreement may be required for certain engagements. To contact the team, email gsfc-dtn@nasa.onmicrosoft.com. Additional information is available on the DTN website: <https://www.nasa.gov/communicating-with-missions/delay-disruption-tolerant-networking/>.

1.4. Assumptions

The guide assumes that node operators are already familiar with BP and DTN, as well as with [cFS Framework](#) and how it integrates with [COSMOS](#) - the chosen telemetry and command system for interfacing with BPNode.

Chapter 2. Software Overview

BPNode is a DTN node that implements BP as defined in RFC-9171. The node software runs on Linux and real-time operating systems in a context where user applications communicate directly with BPNode. BPNode expects a Publish/Subscribe pattern where user applications perform as follows:

- publish Application Data Units (ADUs) to which BPNode subscribes
- subscribe to expected ADUs that BPNode publishes.

The entire system is composed of several modular components grouped by function. They include the core framework, mission-specific applications, supporting libraries, and development and integration tools.

2.1. Supported Platforms

The BPNode application has been built and tested on Ubuntu 20.04.6 Long-Term Support and on Ubuntu 22.04.

2.2. cFS Framework

cFS is a generic flight software architecture framework used on NASA spacecraft and cubesats. The cFS Framework is a core subset of cFS that offers an extensive ecosystem of applications and tools. The [Figure 1, “Framework Structure”](#) figure below shows the external cFS and Core Flight Executive (cFE) framework structure. BPNode interacts with these elements to receive and distribute data for DTN service:

Software Bus

passes the data to BPNode on the application layer. Applications use it to communicate with BPNode and to exchange telemetry, directives, and events.

Event Logging

interacts with **cFE Event Services (EVS)** to collect events generated by BPNode and distribute them via Telemetry Distribution.

Startup

starting, stopping, and restarting the BPNode application is handled by cFE Executive Services (ES)

Time Services

interacts with **cFE TIME** to provide BPNode with time (host time, clock state, etc.) to calculate and maintain the time within the node.

Command Source

interacts with **cFS Command Ingest (CI)** or **Stored Command (SC)** to provide directives to BPNode, such as stored commands or commands received from an external source for BPNode execution.

Telemetry Destination

interacts with **cFS Telemetry Output (TO)** to collect telemetry from BPNode and send it to external entities.

Application

sends and receives Application Data Units (ADUs) to/from BPNode via the framework software bus.

Operating System Abstraction Layer (OSAL)

provides features on the operating system level, including persistent storage and mutexes.

Performance Services

interacts with **cFE Performance Log** to provide logging functionality for BPNode by sending entries through the framework.

Table Services

interacts with **cFE Table Services** to provide a mechanism for table updates which are used by BPNode for configuration purposes.




Figure 1. Framework Structure



See [Chapter 6, Directives, Telemetry, Tables, and Events](#) for the list of BPNode directives, tables, and telemetry.

2.3. BPNode and BPLib

BPNode and BPLib are complementary components that run on top of the cFS framework:

BPLib

serves as the core implementation library that provides BP functionality: <https://github.com/nasa/bplib>

BPNode

integrates this library into the cFS framework and implements the BP node behavior: <https://github.com/nasa/bp>

Together they provide a shared library and support BP and related functions.

BPNode contains the following subsystems, modules, and components:

Application Agent (AA)

provides application and administrative services associated with monitoring and control, such as status monitoring or configuration control:

- Admin Statistics
- Administrative Record Processor
- Framework Proxy
- Node Configuration

BP Agent (BPA)

executes BP-related activities and functions:

- Bundle Interface
- Custody Transfer
- Extension Block Processor
- Payload Interface
- Policy Database
- Storage

Core Infrastructure (CI)

functions as a common utility across modules for low-level functions and processing, such as memory, time, and events:

- Concise Binary Object Representation (CBOR)
- Cyclic Redundancy Check (CRC)

- Endpoint ID (Endpoint)
- Event Management
- Memory Allocator
- Performance Logger
- Queue Management
- Red-Black Tree
- Time Management

Convergence Layer Adapters (CLAs)

interfaces with external Convergence Layers (CLs) such as UDP to send and receive bundles on behalf of the BPA



The following features only contain skeleton code in build 7.0.5 and are slated for future implementation: Policy Database

2.4. Dependencies

For both BPNode and BPLib, the following dependencies are required if running on Ubuntu 20.04.06:

CMake (build system)

version 3.22.1 used during testing

pkg-config

version 0.29.1 used during testing

Gnu Compiler Collection toolchain

version 9.4.0 used during testing

sqlite3

version 3.31.1 used during testing

BPNode and BPLib have also been tested on Ubuntu 22.04 with the following package versions:

CMake (build system)

version 3.22.1 used during testing

pkg-config

version 0.29.2 used during testing

Gnu Compiler Collection toolchain

version 11.4.0 used during testing

sqlite3

version 3.37.2 used during testing

These packages can be installed as follows:

```
sudo apt-get install cmake pkg-config build-essential libsqlite3-dev
```

2.5. QCBOR

BPLib relies on CBOR to handle the serialization and deserialization of data structures transmitted over the protocol. BPLib uses version 1.5.1 of QCBOR for CBOR data encoding and decoding.

Follow these steps to install QCBOR:

1. Clone the QCBOR repository:

```
cd ~/ # Or any directory for installing QCBOR
git clone https://github.com/laurencelundblade/QCBOR.git
```

2. Navigate to the QCBOR repository:

```
cd QCBOR
```

3. Select the correct version:

```
git checkout v1.5.1
```

4. Build and install the library:

```
sudo make uninstall # Remove any previous installation
cmake -DBUILD_SHARED_LIBS=ON -S . -B build
cmake --build build
sudo make install
```

Chapter 3. Installation

3.1. Building BPNode and BPLib within cFS

3.1.1. Clone Repositories

If you do not already have a cFS bundle set up, clone the cFS repository to the local system:

```
cd <chosen working directory>
git clone https://github.com/nasa/cFS.git
cd cFS
export CFS_HOME="$(pwd)" # Set the home directory for cFS
```

Follow the instructions in the cFS README to initialize the sample cFS deployment: <https://github.com/nasa/cfs?tab=readme-ov-file#build-and-run>



cFS relies heavily on Git submodules, so to avoid errors when updating code, they must be kept updated by running `git submodule update --init`.

Clone the BPNode and BPLib repos as submodules in the cFS bundle.

```
cd "${CFS_HOME}"/apps
git submodule add https://github.com/nasa/bp.git bpnode
cd "${CFS_HOME}"/libs
git submodule add https://github.com/nasa/bplib.git
```

The CF app, provided by cFS, is used by the sample configurations provided as a source of bundle ADUs. To add it to the build system, clone it as a submodule in your cFS bundle:

```
cd "${CFS_HOME}"/apps
git submodule add https://github.com/nasa/CF.git cf
```

The DTN test tools are a separate python toolkit that does not depend on cFS but can be used for DTN testing either in COSMOS or on the command line. This repository can be cloned into the directory of your choice, it will not affect the cFS build system.

```
git clone https://github.com/nasa/dtn-tools.git
```

3.1.2. Adding BPNode Configurations to cFS

The BP repository includes a definitions directory named `dtn_defs` with all the configurations needed to instantiate a sample multi-node DTN configuration with custody transfer. If you aren't starting a new cFS configuration with DTN, you can simply copy relevant parts of the files into your

configurations. Otherwise, copy the whole directory into your top-level cFS directory.



If you followed the cFS instructions when instantiating your cFS bundle, remove the sample_defs directory, the dtn_defs directory replaces it. Please make sure the Makefile provided by cFS is still included at the top-level cFS directory.

```
mv "${CFS_HOME}"/apps/bpnode/dtn_defs "${CFS_HOME}"
```

3.1.2.1. dtn_defs Content

For more details on the BPNode tables, please refer to the provided DTN M&C ICD. Special items of note in creating a 2-node custodial transfer with BPNode tables are called out here.

dtn_defs/config/

cf_topicids.h

This file defines the topic IDs to values expected by our COSMOS plugin

dtn_defs/cpu<x>/tables/

bpnode_adup.c

This version of the ADU proxy table is set to use the CFE_SB_STATS, CFE_SB_ALLSUBS, CFE_SB_ONESUB, and CF_PDU packets as ADUs in bundle creation. The CF_PDU message ID in particular can be used to create delay and disruption tolerant file transfers between two nodes.

bpnode_channel.c

The table's block configurations define what extension blocks to add to bundles created on each channel. For instance, the IncludeBlock member of the CustodyTransferBlkConfig being set to true defines all bundles on a given channel as custodial.

bpnode_contacts.c

The CSSizeTrigger and CSTimeTrigger values define the transmission rates of Compressed Custody Signals on each contact. RetransmissionTimeouts define how frequently custodial bundles are retransmitted.

bpnode_mib_pn.c

The PARAM_SUPPORT_CUSTODY member of this table is used to determine if a node supports custody transfer operations. The node's endpoint identifier is also defined here, and is used by other node's contact tables to determine bundle routing.

dtn_defs/cpu<x>/cfe_es_startup.scr

This startup script is modified to have BPLib and BPNode started automatically with cFS.

dnt_defs/tables/

sch_lab_table.c

This Scheduler table is modified to add BPNode wakeup messages at 10Hz and telemetry requests at 0.25Hz. This table is just a recommendation though. The telemetry request rate

can be customized entirely to a project's needs. The wakeup rate is recommended to be set to 10Hz and should not be set to slower than 1Hz. See `BPNODE_MAX_EXP_WAKEUP_RATE` in [bpnode/config/default_bpnode_platform_cfg.h](#) for more details.

to_lab_sub.c

The TO subscription table was modified to ensure BPNode telemetry packets are egressed by `to_lab`.

dtndefs/targets.cmake

BPNode and BPLib have been inserted into the `MISSION_GLOBAL_APPLIST` variable in this modified `targets.cmake` file

dtndefs/mission_build_custom.cmake

`mission_build_custom.cmake` needs to have BPNode's PSP module for UDP socket interfacing (`udpsock_intf`) added to the build system

Configurations are provided to build 5 DTN nodes as 5 separate cFS instances, but only `cpu1` and `cpu2` are configured to communicate with each other out of the box. For additional context on how to build multiple cFS instances as separate CPUs, see [here](#).

3.1.3. Building and Deploying the DTN cFS

This method is ideal for simulation, unit testing, and general development work on a local Linux machine. Use the following commands from the top-level `cFS` directory:

```
make distclean
make SIMULATION=native prep
make
make install
cd build/exe/cpu1/
./core-cpu1
```

This compiles and launches the `cpu1` target in a native environment.

3.1.4. Run Unit Tests and Generate Coverage

Testing is a key part of validating the DTN cFS system. The sequence below starts by cleaning and configuring the native build environment to include unit test support, followed by compilation, testing, and coverage reporting.

```
make distclean
make SIMULATION=native ENABLE_UNIT_TESTS=true prep
make
make test
make lcov
```

The resulting report can be used to assess code coverage and identify gaps in test validation.



Always run `make distclean` before switching targets or simulation environments to avoid build conflicts. For advanced build configurations or debugging, consult the Makefiles and scripts included in the repository.

3.1.5. Troubleshooting

This section addresses common issues that may arise when working with the DTN cFS bundle.

Missing files or directories after cloning

Run `git submodule update --init` to ensure all nested submodules are pulled.

Unit tests not executing

Make sure to include `ENABLE_UNIT_TESTS=true` in the `make prep` step.

3.2. Standalone BPLib Build (bpcat)

BPLib can be built as a standalone executable called `bpcat`, suitable for a ground relay node or as a demo of bplib's features without the full CS architecture. This standalone build still depends on the OSAL library.



Note that `bpcat` is only a prototype and does not implement the full suite of BPNode features, since the node relies on cFS for directive handling, telemetry routing, table updates, etc.

Build/run instructions for `bpcat`:

1. Clone the repositories

```
cd <chosen working directory>
export BPCAT_HOME="$(pwd)" # Set working directory
git clone https://github.com/nasa/bplib "${BPCAT_HOME}"/bplib
git clone https://github.com/nasa/osal "${BPCAT_HOME}"/bplib/osal
```

2. Configure the build environment

```
# Choose a build configuration - Debug or Release
export MATRIX_BUILD_TYPE=Debug

# Indicate where the OSAL build files are
export NasaOsai_DIR="${BPCAT_HOME}/bplib/osal-staging/usr/local/lib/cmake"

# Define bplib build directory
export BPLIB_BUILD="${BPCAT_HOME}/bplib/bplib-build-matrix-${MATRIX_BUILD_TYPE}-POSIX"
```


3. Build BPLib with OSAL

```
cd $BPCAT_HOME/bplib/osal
cmake -DCMAKE_INSTALL_PREFIX=/usr/local \
      -DOSAL_SYSTEM_BSPTYPE=generic-linux \
      -DCMAKE_BUILD_TYPE="${MATRIX_BUILD_TYPE}" \
      -DOSAL_OMIT_DEPRECATED=TRUE \
      -DENABLE_UNIT_TESTS=TRUE \
      -DOSAL_CONFIG_DEBUG_PERMISSIVE_MODE=ON \
      -B ../osal-build
cd ../osal-build
make DESTDIR=../osal-staging install
cd ..
cmake -DCMAKE_VERBOSE_MAKEFILE=ON \
      -DCMAKE_BUILD_TYPE="${MATRIX_BUILD_TYPE}" \
      -DBPLIB_OS_LAYER=OSAL \
      -DCMAKE_PREFIX_PATH=/usr/local/lib/cmake \
      -B "${BPLIB_BUILD}"
cd $BPLIB_BUILD
make all
```

4. Test and execute

- To run the bplib unit tests:

```
cd $BPLIB_BUILD
ctest --output-on-failure 2>&1 | tee ctest.log
```

- To run bpcat as a standalone:

```
cd $BPLIB_BUILD
cd app
./bpcat
```

In the default bpcat configurations, a single contact is started that listens on 0.0.0.0:4501 and forwards bundles with the EID ipn:200.64 to 127.0.0.1:4551. To change these configurations, modify the ContactsTbl in bplib/app/src/bpcat_nc.c and rebuild bplib.

Chapter 4. Configuration

4.1. Startup Configuration

On startup, the BPNode main task will initialize itself, the BPLib library data, and any child tasks. On first initialization, the node creates a SQLite database named `bplib-storage.db` in the same directory as the executable (typically `build/exe/cpu1/`) along with temporary files associated with it. On subsequent initializations BPLib reads from it and reports the total number of bundles stored (see the Node MIB Reports telemetry packet as described in [Chapter 6, Directives, Telemetry, Tables, and Events](#)).

The node's Endpoint Identifier (EID) is established during initialization by reading from the node MIB configuration table. This table contains an "instance EID" field that provides the base EID value used as the node's primary identity. Channels use this EID for bundles originating on local applications, modifying the service number to match the corresponding value from the channel configuration table.



See [Chapter 6, Directives, Telemetry, Tables, and Events](#) for more information.

The counts of tasks are determined by the following BPNode or BPLib macros:

BPNODE_NUM_GEN_WRKR_TASKS

Generic Worker

BPLIB_MAX_NUM_CHANNELS

ADU In/Out

BPLIB_MAX_NUM_CONTACTS

CLA In/Out

There is also one maintenance task whose count is not configurable.

4.1.1. ADU Channels

ADU tasks are implemented via virtual communication pathways for application data called channels. Each channel has a unique ID from 0 to `BPLIB_MAX_NUM_CHANNELS` - 1. All channels default to `REMOVED` state at initialization unless their configuration in the channel configuration table has `loadAutomatically` set to `true`, in which case they should initialize as `STARTED`. Channel ID indexes into:

- The ADU proxy table
 - This table has any cFS-specific channel information, including which message IDs a channel should subscribe to
- The channel configuration table
 - This table has general channel configurations, including the service number to assign to the channel, and how to configure new bundles created from this channel's ADUs

- The channel status array in the channel/contact status telemetry packet



See [Chapter 6, Directives, Telemetry, Tables, and Events](#) for more information on BPNODE tables.

4.1.2. CLA Contacts

CLA tasks are implemented via network interface connections called contacts: Each contact has a unique ID from 0 to `BPLIB_MAX_NUM_CONTACTS` - 1. All contacts default to `TORNDOWN` state at initialization. Contact ID indexes into:

- The contact configuration table
 - This table has information including which destination EID(s) to associate with this contact and the needed IP information to set up a UDP contact
- The contact status array in the channel/contact status telemetry packet

Chapter 5. Running the Node

Main configurations that drive node performance:

Rate limits

see the channel and contact configuration tables

Task priorities

see `bnode/config/default_bnode_platform_cfg.h` and the cFE startup script

Memory pool size

see `BPNODE_MEM_POOL_LEN` in `bnode/config/default_bnode_platform_cfg.h`

Number of tasks in the system

based on `BPLIB_MAX_NUM_CONTACTS`, `BPLIB_MAX_NUM_CHANNELS`, and
`BPNODE_NUM_GEN_WKR_TASKS` in `bplib/inc/bplib_cfg.h` and
`bnode/config/default_bnode_platform_cfg.h`

The maximum bundle bytes allowed in the storage database

based on `BPLIB_MAX_STORED_BUNDLE_BYTES` in `bplib/inc/bplib_cfg.h`

The maximum size of the Custodial Transfer Database

based on `BPLIB_CT_DB_MAX_ENTRIES` in `bplib/inc/bplib_cfg.h`

Whether duplicate bundles are allowed in storage

see `BPLIB_ALLOW_DUPLICATE_BUNDLES` in `bplib/inc/bplib_cfg.h`

The batch size of bundles

loaded out of storage

`BPLIB_STOR_LOADBATCHSIZE`

loaded into storage

`BPLIB_STOR_INSERTBATCHSIZE`

discarded from storage

`BPLIB_STOR_DISCARDBATCHSIZE`

Note that configurations can affect each other. The storage bundle byte limit and the memory pool size should be at least as large as the average expected bundle size times the relevant load batch sizes into or out of storage.

Configurations that affect how bundles move between nodes and how bundles are created/processed are primarily found in the ADU proxy, channel configuration, and contact configuration tables. Multiple BPNode instances can be deployed in a network as source, relay, or destination nodes, as shown in [Figure 2, “BPNode deployment roles”](#) below:



Tables located in the top-level cFS `[MISSION]_defs` folder will always override the table definitions in `apps/bnode/fsw/tables/`.




Figure 2. BPNode deployment roles

5.1. Managing Channel Applications

BPNode manages channel applications by separating configuration from activation, allowing for complete lifecycle control from setup through teardown. When an application needs to be integrated into the network, it must first be added to make its configuration available, then explicitly started to begin data flow. Applications can later be stopped to halt data processing or completely removed, which flushes any remaining data from both the egress queue and the cFE software bus pipe. These application lifecycle operations are controlled through the following directives:

add-application

this directive loads an application's configuration by reading settings from both the ADU proxy table and channel configuration table using the specified channel ID. However, this step only makes the configuration available - no actual data processing begins until the application is started.

start-application

this directive activates the configured application by enabling its ADU tasks to begin processing data flow. In the cFS environment, this activation involves subscribing to the relevant ADU message IDs and establishing the connection to post delivered bundle payloads to SB. Note that the registration state of a channel will determine whether ADUs are delivered, stored, or discarded.

stop-application

this directive stops the flow of ADUs in and out of the specified channel.

remove-application

this directive flushes the egress queue of the specified channel and stores any bundles found in there. It will also flush the cFE SB pipe of any new ADUs.



See [Chapter 6, Directives, Telemetry, Tables, and Events](#) for the list of BPNode directives.

Once active, applications participate in BP communication through two primary data flows:

ADU In

this task handles data originating from local applications by converting received cFS packets into bundle payloads, applying the settings specified in the channel configuration table to create properly formatted bundles for network transmission.

ADU Out

this task handles inbound data bundles arriving from the network that are destined for local applications. When a bundle's destination EID matches an application running on the current node, the bundle is routed to the appropriate task, which extracts the payload data and posts it to the cFS SB where the target application can receive it.

This architecture allows BPNode to serve as a bridge between local cFS applications and the broader BP network, handling the translation between application-specific data formats and standardized bundle protocols.

5.2. Managing Contacts

BPNode manages network connections through contacts using a comprehensive lifecycle process similar to application management, allowing for complete control from initial setup through final teardown. Contacts represent network connections to other BP nodes and must be configured before they can actively participate in bundle routing. Like applications, contacts can be stopped to halt bundle traffic or completely torn down, which flushes any remaining bundles from the egress queue into storage. These contact lifecycle operations are controlled through the following directives:

contact-setup

this directive establishes a contact's network configuration by reading settings from the contact configuration table using the specified contact ID. This step creates the underlying CL connection infrastructure but leaves it in a dormant state with no bundle traffic flows until the contact is explicitly started.

contact-start

this directive activates the configured contact by enabling its CLA tasks to begin processing bundle traffic. Once started, the contact can participate in both receiving bundles from and transmitting bundles to other nodes in the BP network.

contact-stop

this directive stops the flow of bundles in and out of the specified contact.

contact-teardown

this directive flushes the egress queue of the specified contact and stores any bundles found in there.



See [Chapter 6, Directives, Telemetry, Tables, and Events](#) for the list of BPNode directives.

Active contacts handle BP communication through two distinct pathways:

CLA In

this task receives encoded bundles from incoming network traffic, decodes the CBOR formatting, and deserialize the data into a readable bundle format. Each received bundle undergoes validation to ensure data integrity. Bundles that fail validation are immediately discarded to prevent corrupt data from entering the system.

CLA Out

this task manages the outgoing traffic by routing bundles destined for remote nodes based on their destination EIDs. When a bundle's destination corresponds to a destination EID in the started contact, it is forwarded to the appropriate task, which encodes the bundle using CBOR formatting and transmits it to the next node in the network routing path.

This contact management system enables BPNode to maintain reliable connections with multiple network peers while ensuring data integrity through validation and proper encoding protocols.

5.3. Bundle Storage Management

When BPNode receives a bundle that cannot be immediately processed—either because its destination EID does not match any active local applications (channels) or available network connections (contacts), the system provides persistent storage to ensure reliable message delivery even across system restarts. Such bundles are written to the SQLite database `bplib-storage.db`, which provides persistent storage that survives node restarts, ensuring that bundles waiting for delivery or forwarding opportunities are not lost during system maintenance or unexpected shutdowns. Custodial bundles, which will be retransmitted if reception by the next node is not confirmed in time, are always stored, except when immediate delivery to the destination application is available.

The main task continuously monitors stored bundles to manage their lifecycle effectively. It regularly checks each bundle's creation timestamp against its specified lifetime to identify expired bundles, which are automatically discarded to prevent indefinite storage of undeliverable messages. Simultaneously, the system monitors for new routing opportunities by checking whether previously unroutable bundles can now be forwarded to newly available contacts or delivered to recently started local applications. When such opportunities arise, bundles are moved from storage to the appropriate egress queue and removed from the database.

The system provides telemetry reporting on storage utilization through the storage telemetry packet, which includes both the total size of the `bplib-storage.db` file and the actual space consumed by stored bundles. The bundle storage size is typically smaller than the total database file size due to SQLite's internal database management overhead. Storage capacity is controlled by the `BPLIB_MAX_STORED_BUNDLE_BYTES` macro, which sets limits based on the actual byte size of bundle content rather than total database file size.

Detailed developer notes on the storage design can be found in the BPLib docs folder: <https://github.com/nasa/bplib/tree/main/docs>

5.4. Running a Sample Bundle Flow

Assuming you've made and installed cFS as described in [Section 3.1.3, "Building and Deploying the DTN cFS"](#), you should have five compiled cFS instances in `build/exe/`. This demo will use two instances, one to demonstrate a source node and one to demonstrate a destination node. Open two terminals, one for each instance.

In the first terminal, run the following:

```
cd build/exe/cpu1/  
./core-cpu1
```

This terminal represents your source node.

In the second terminal, run the following:

```
cd build/exe/cpu2/  
./core-cpu2
```

This terminal represents your destination node.

Follow [Section 9.5, “COSMOS Integration”](#) to set up your COSMOS instance. Familiarity with COSMOS is assumed. These instructions assume you are running COSMOS from the same machine as the two DTN nodes, if that is not the case, replace any IP addresses specified here with the relevant ones.

In COSMOS, send the following commands to set up your nodes:

```
cmd("DTNFSW-1 TO_LAB_CMD_ENABLE_OUTPUT with DEST_IP 127.0.0.1")  
cmd("DTNFSW-2 TO_LAB_CMD_ENABLE_OUTPUT with DEST_IP 127.0.0.1")  
  
cmd("DTNFSW-1 CFE_TIME_CMD_SET_STATE with CLOCK_STATE 'VALID'")  
cmd("DTNFSW-2 CFE_TIME_CMD_SET_STATE with CLOCK_STATE 'VALID'")
```

DTN time in bundles only gets set if the host system is reporting its time data as valid. The DTN timestamps will likely not be accurate to the real-world time, but configuring cFE Time Services is out of the scope of this user guide.

Then, set up your contacts:

```
cmd("DTNFSW-1 BPNODE_CMD_CONTACT_SETUP with CONTACT_ID 0")  
cmd("DTNFSW-1 BPNODE_CMD_CONTACT_START with CONTACT_ID 0")  
  
cmd("DTNFSW-2 BPNODE_CMD_CONTACT_SETUP with CONTACT_ID 0")  
cmd("DTNFSW-2 BPNODE_CMD_CONTACT_START with CONTACT_ID 0")
```

This is the contact both nodes will be using to communicate with each other. Node 1 (core-cpu1) will be creating bundles on channel 0, which is started automatically at initialization, and node 2 (core-cpu2) will deliver bundles to channel 0, which is also started automatically.

To transfer a file between the two nodes, copy a file of your choosing into build/exe/cpu1/cf/. For the purposes of this exercise, we will assume the file is called sample.txt.

Initiate a custodial bundle transfer between node 1 and node 2 with the following CF command:


```
cmd("DTNFSW-1 CF_CMD_TX_FILE with CFDP_CLASS 'CLASS_1', KEEP 1, CHAN_NUM 0, PRIORITY  
0, DEST_ID 2, SRC_FILENAME '/cf/sample.txt', DST_FILENAME '/cf/sample.txt'")
```

In the BPNode COSMOS telemetry screens, you should see various bundle counters on node 1 and node 2 increment. When the file transfer is complete, you should see the file copied to build/exe/cpu2/cf/.

Chapter 6. Directives, Telemetry, Tables, and Events

For the list and descriptions of BPNode directives, telemetry, tables, and events, consult the Monitor & Control Interface Control Document (M&C ICD) in the BPNode repository.

Chapter 7. Custody Transfer

BPNode follows the Custody Transfer specification laid out in the CCSDS Compressed Bundle Status Reporting and Custody Signaling Orange Book draft (CCSDS 734.6-O-1). Since the Orange Book is currently in draft, the details of how Custody Transfer is implemented in BPNode are described in detail in this section.

Custody Transfer enables acknowledgment of bundle transfer between DTN nodes by shifting storage and retransmission responsibility from one node to the next. When a node accepts custody of a bundle, it becomes responsible for storing and retransmitting that bundle until it receives an accepted custody signal or the bundle expires. If the bundle is on its destination node, custody is automatically accepted and confirmed by the standard ADU delivery process, instead of through a custody signal. Bundles request custody by including a Custody Transfer Extension Block (CTEB), and custody decisions are communicated through Compressed Custody Signals (CCS), a type of administrative record bundle, that can acknowledge multiple bundles efficiently using sequence number ranges. The custody transfer process is integrated with contact lifecycle management, with CCS size and time trigger values configured during contact-setup and active timers cancelled upon contact-teardown.

The Custody Transfer (CT) Module:

- maintains the Custody Transfer Database (CTDB) that maps bundle IDs to a sequence ID and sequence number pair
- updates a bundle's CTEB when forwarding the bundle to a new node
- processes received CCS bundles and updates the CTDB accordingly
- manages in-progress CCS information per contact
- handles size and time triggers for CCS transmission
- manages retransmission timers and cancellation when custody is rejected or when a contact is torn down
- makes custody acceptance/rejection decisions based on storage availability and policy authorization

7.1. Bundle Definitions

Bundles are marked as custodial by the inclusion of a Custody Transfer Extension Block (CTEB), block type 13. The block type-specific data is:

Table 1. CTEB Definition

Field Name	Field Type	Description
Sequence number	unsigned integer	With the sequence ID, a unique identifier for this bundle assigned by the previous node that had custody of this bundle

Sequence ID	unsigned integer	With the sequence number, a unique identifier for this bundle assigned by the previous node that had custody of this bundle
Block Source Administrative EID	EID array	The EID of the previous node that had custody of this bundle

Custodial acknowledgement of bundles is handled by the CCS bundle, administrative record type 13. The record content is a CBOR mapping of disposition codes to bundle sequence collections. Disposition codes are signed integers that, if positive, indicate that the associated bundle sequence collection was accepted, and if negative, indicate that the collection was rejected. Bundle sequence collections are CBOR arrays as follows:

Table 2. CCS Bundle Sequence Collection (BSC)

Field Name	Field Type	Description
Sequence ID	unsigned integer	Unique identifier for the counter that the sequence numbers in this BSC belong to. Each new contact that is started gets assigned a new sequence ID starting at 1.
First sequence number	unsigned integer	Starting point for the sequence range lengths
Sequence range lengths	array	An array with an odd number of unsigned integers that specify, alternately, ranges of included and excluded bundle sub-sequences

7.2. Data Flow

The custody transfer mechanism involves two distinct but complementary flows that handle different directions of the custody process. The first flow covers how a node processes outgoing bundles when it needs to request custody from a downstream node, while the second flow handles incoming custody acknowledgments that complete or reject previous custody requests from an upstream node.

When a bundle is processed for custody transfer, the system follows a careful sequence to ensure reliability. The bundle cannot be accepted for custody until it's safely stored, preventing data loss if the node fails during processing. The complete custody transfer process spans multiple nodes and includes decision points for storage availability, authorization, and policy compliance.




Figure 3. Custodial Bundle Handling Executed - End-to-End Flow

7.2.1. When to Accept or Reject Custody

If a bundle contains a CTEB, the node evaluates whether to accept or reject custody of it following a systematic evaluation of storage and authorization constraints. Custody acceptance is not considered complete until the bundle has been successfully stored.




Figure 4. Custody Decision and Bundle Storage Process

The current iteration of BPNode follows these criteria when determining whether to accept or reject custody of a custodial bundle:

1. Custody is rejected and the bundle is discarded if:
 - a. There is not enough storage available
 - b. The CTDB cannot hold more entries
 - c. The bundle is destined for a local channel but that channel's registration state is set to passive-abandon.
 - d. The storage operation failed for some unknown reason
2. Custody is accepted but the bundle is discarded if the bundle is a duplicate of a bundle already in the CTDB (retransmission detected)
3. Otherwise custody is accepted, the bundle is stored, and the bundle is added to the CTDB.

Future versions of BPNode will implement the Policy Database to allow for more comprehensive network policies governing the acceptance of custodial nodes.

7.2.2. Custodial Bundle Handling

Whenever a custodial bundle is received, its sequence ID and sequence number is added to an in progress CCS to signal its acceptance or rejection to the previous node.

Once a custodial bundle is accepted and stored, storage will attempt to egress it on the first applicable contact. When a contact is started, it is assigned a new sequence ID with an associated monotonically increasing sequence counter. This determines the sequence ID information to use for custodial bundles leaving the node over this contact. When the bundle leaves the node for the

first time, the bundle's sequence ID and sequence number are assigned, added to the CTDB for tracking, and the bundle's CTEB is updated with the new values. Until a CCS is received acknowledging custody of this bundle's sequence ID/number pair, that bundle will be periodically retransmitted until it expires.

7.2.3. Custody Signal Processing Flow

The custody signal processing flow handles the reverse direction - when a node receives acknowledgments for bundles it previously sent with custody requests. This process is crucial for completing the custody transfer and freeing up storage resources. Custody signals are generated when any of the following trigger conditions are met:

1. Time trigger: Defined in the contact table, when a certain amount of time elapses since bundles began being added to a CCS in memory
2. Size trigger: Defined in the contact table, when the CCS is estimated to have reached a certain size
3. CCS count trigger: Defined by the `BPLIB_CT_MAX_OPEN_CCS` macro, when the maximum number of in progress CCSs in memory is reached, the oldest CCS is sent
4. Sequence range length trigger: Defined by the `BPLIB_CT_MAX_SEQ_RANGE_LEN` macro, when the maximum bundle sequence range length is reached
5. contact-stop trigger: A contact's CCS are all generated when that contact is stopped




Figure 5. Custody Signal Creation and Transmission Process

When a Compressed Custody Signal arrives, it contains acknowledgments (acceptances or rejections) for multiple bundles encoded as sequence number ranges. The system first validates the signal format and source, then processes each range to determine which bundles were accepted or rejected by the destination node. Missing bundles are also implicitly acknowledged by noting the gaps in sequence numbers.

For each acknowledged bundle, the system looks up the corresponding bundle in its CTDB using the sequence number and ID.

1. If the bundle was accepted by the destination node, this completes the custody transfer and the node can delete its stored copy.
2. If the bundle was rejected, the node turns off bundle retransmission but it retains custody responsibility and may resume retransmission if a new contact is started.
3. If the bundle was missing, the node automatically triggers bundle retransmission.

7.3. Configuration

The following configurations are needed to enable custody transfer flows in a node:

1. Channel Table
 - a. Set the CTEB configuration's inclusion flag to true to make bundles generated on this channel custodial
2. Contact Table
 - a. Size Trigger: Maximum size an in-progress CCS can be before it is generated (measured in bytes of the predicted encoded bundle)
 - b. Time Trigger: Maximum time an in-progress CCS can sit in memory before it is generated
 - c. Retransmission Timeout: Maximum time to wait for custody acknowledgement before retransmitting bundles sent on this contact
3. Node MIB Configuration Table
 - a. Set the field at index `PARAM_SUPPORT_CUSTODY` (4) to true

7.4. Operation

Custody transfer flows are controlled like all other bundle flows with contact and channel directives. The only difference is that contact-setup will update retransmission triggers for that contact in the node, while contact-teardown will disable them. Additionally, contact-stop will trigger the generation of all CCSs relevant to that contact.

To track custody transfers in the system, the following MIB configurations, counters and reports are called out:

Table 3. Key Performance Counters

Field Name	Packet Name	Description
<code>paramSupportCustody</code>	Node MIB Config	Whether this node can process custody transfers
<code>bundleCountInCustody</code>	Node MIB Reports	Bundles currently in custody of this node
<code>BundleAgentCtdbSize</code>	Node MIB Reports	Maximum allowed size of the CTDB
<code>bundleCountCustodyRequest</code>	Node MIB Counters	Bundles that requested custody from this node
<code>bundleCountAcceptedCustody</code>	Node MIB Counters	Bundles that this node accepted custody of
<code>bundleCountRejectedCustody</code>	Node MIB Counters	Bundles that this node rejected custody of
<code>bundleCountCustodyTransferred</code>	Node MIB Counters	Bundles whose custody was successfully transferred to another node

bundleCountCustodyRejected	Node MIB Counters	Bundles whose custody was rejected by the next node
bundleCountCustodyReForwarded	Node MIB Counters	Number of bundle retransmissions
bundleCountGeneratedCustodySignal	Node MIB Counters	Bundles that this node generated a custody signal for in a CCS
bundleCountReceivedCustodySignal	Node MIB Counters	Bundles for which custody signals were received by this node in a CCS
bundleCountCcsReceived	Node MIB Counters	Compressed Custody Signal bundles received by this node
bundleCountGeneratedCcs	Node MIB Counters	Compressed Custody Signal bundles generated by this node

Chapter 8. Known Issues

1. If the general rate of ingressing data is greater than the general rate of egressing data, memory will start to run out and ingressing bundles will start to be dropped.
2. If bundles start to expire from storage at the same time as they start to be egressed from storage, the `BUNDLE_COUNT_LIFETIME_EXPIRED` telemetry counter might erroneously count bundles that have been egressed as expired.

For a full list of issues, see the BPNode 7.0.5 Version Description Document. These issues are all slated for fixing in build 7.1.

Chapter 9. DTN Tools Suite

The [DTN Tools suite](#) is a toolkit that allows for:

- confirming whether bundles created by the implementation adhere to specifications such as RFC-9171, ensuring protocol compliance at a fundamental level
- rigorous testing of how the implementation handles invalid or off-nominal bundles
- verifying implementation's resistance to corrupted data

The toolkit provides comprehensive end-to-end tests within a single script for a consistent methodology for creating bundle content, covering the entire process from bundle creation through transmission to a DTN node, reception from a DTN node, and verification of received bundles.




Figure 6. Tools Overview

It includes support for extension blocks like Custody Transfer Extension Block (CTEB) and Compressed Reporting Extension Block (CREB) as well as CBOR-encoded and decoded formats.

The suite contains two tools: DTNGEN and DTNCLA.

9.1. DTNGEN

DTNGEN is a Python library that enables the creation and interpretation of DTN bundles for compliance verification as well as error handling and interoperability testing between different DTN implementations. It can create both standards-compliant bundles and deliberately invalid bundles with incorrect data or block ordering.

The tool provides bundle manipulation capabilities through dedicated classes for primary blocks and canonical block types. It can generate:

- individual bundles with customizable components (primary, payload, CCS administrative record content, and extension blocks)
- sets of bundles using:
 - script-based loops
 - specialized Bundle Set Generation feature, which allows for controlled variations in creation

timestamps and payload sizes.




Figure 7. Single Bundle and Bundle Set Generation

DTNGEN offers flexible data format support, allowing bundles to be read from and written to both binary and JSON formats:




Figure 8. Conversion between CBOR and JSON

```

{
  "block_class": "PrimaryBlock",
  "version": 8,
  "control_flags": 4,
  "crc_type": 1,
  "dest_eid": {
    "type": "EID",
    "value": {
      "uri": 2,
      "ssp": {
        "node_num": 103,
        "service_num": 1
      }
    }
  },
  "src_eid": {
    "type": "EID",
    "value": {
      "uri": 2,
      "ssp": {
        "node_num": 101,
        "service_num": 1
      }
    }
  },
  "rpt_eid": {
    "type": "EID",
    "value": {
      "uri": 2,
      "ssp": {
        "node_num": 100,
        "service_num": 1
      }
    }
  },
  "creation_timestamp": {
    "type": "CreationTimestamp",
    "value": {
      "time": 755533883904,
      "sequence": 0
    }
  },
  "lifetime": 3600000,
  "crc": {
    "type": "hexbytes",
    "value": "af7c"
  }
}

```

9.2. DTNCLA

DTNCLA is a unified Python library that implements CLAs needed for DTN testing. It provides a consistent interface for communicating with DTN nodes, regardless of the underlying convergence layer protocol.

The library follows a modular design where most functionality is shared across different CLAs, promoting code reuse and consistent behavior. Only the network read/write functions are specialized, with implementations tailored to the specific requirements of each convergence layer protocol.

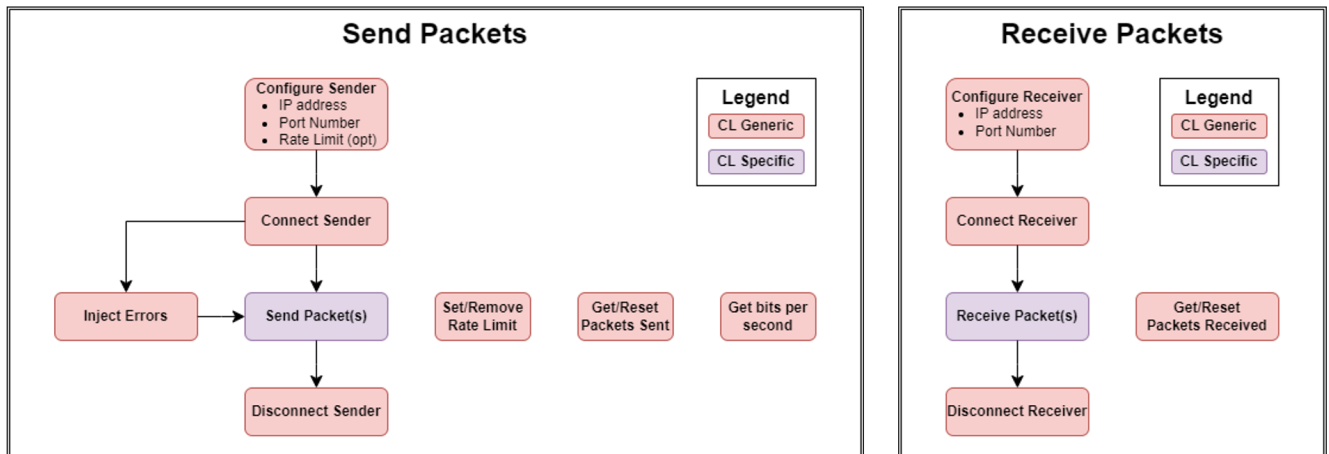


Figure 9. Send and Receive Packets

9.3. Dependencies

The suite requires the following dependencies:

- Python v3.8 or higher
- cbor2 v5.6.2 or higher
- crccheck v1.3.0 or higher

9.4. Tools Suite Installation

Each DTN Tools Suite release provides a complete package containing a built Python wheel file (.whl) for easy installation, source code, release notes, and documentation in HTML and PDF formats that covers all classes, variables, and usage examples that can be copied directly into test scripts.

To install for command line usage:

```
pip install dtntools-X.X.X-py3-none-any.whl
```

The source code is organized into four main folders:

dtntools

Contains the Python source code for DTNGEN and DTNCLA

docs

Houses all HTML documentation

tests

Contains scripts for testing DTN Tools, including unit tests and tests against existing DTN implementations

examples

Includes demonstration scripts for DTN Node testing:

bundle-generation

Shows how to generate bundles, send them to a DTN Node, receive them back, and verify content

multiple-contacts

Demonstrates using Python "subprocess" to simulate multiple concurrent contacts

performance-test

Illustrates creating and exchanging large bundle sets with a DTN node at configured bit rates

9.5. COSMOS Integration

The DTN Tools Suite has verified integration with OpenC3 COSMOS versions 5.17.1 and 6.3.0, which support native Python package installation. DTN Tools packages can be installed directly through the COSMOS Administrator Console, with required dependencies (cbor2 and crccheck) being automatically resolved and installed during the process.

When configuring test environments, any host ports or folders accessed by the DTN test scripts must be explicitly defined in the COSMOS compose.yaml file to ensure proper container communication. The integration enables calling DTN Tools functions directly from COSMOS script-runner, facilitating their incorporation into scripts and procedures.

```
96  try:
97      data_receiver.connect()
98      data_sender.connect()
99
100     for x in (bundle_data):
101         data_sender.write(x)
102
103     for x in range(1, len(bundle_data)+1):
104         received_bundle = data_receiver.read()
105         decoded_bundle = Bundle.from_bytes(received_bundle)
106         if decoded_bundle:
107             print(decoded_bundle.to_json())
108             decoded_bundle.to_json_file(f"/bundles/received_bundle_{x}.json")
109
110     print(f'Packets sent = {data_sender.get_packets_sent()}')
111     print(f'Packets received = {data_receiver.get_packets_received()}')
112
```

Figure 10. Bundle transmission and reception integrated into COSMOS script-runner



Within the COSMOS environment, some methods differ from standard Python constructs, e.g.: COSMOS uses "ask" rather than "input" for user prompts.

Appendix A: COSMOS Interface

COSMOS is a suite of software applications for integration, testing, and operations of embedded systems that can range from power supplies to satellite constellations. BPNode has been tested against OpenC3 COSMOS versions 5.17.1 and 6.3.0. All files needed to run BPNode against COSMOS are provided in the `bpnode/openc3-cosmos-dtnfsw` directory.

To run COSMOS with the `openc3-cosmos-dtnfsw` plugin, first make the following modifications to COSMOS's `compose.yaml` file:

- If you are running `to_lab`, add `"1235:1235/udp"` to the ports list under the `openc3-operator` section

```
openc3-operator:
  ports:
    - "1235:1235/udp"
```

- If you are using one or more UDP configurations for the CLA tasks and wish to use the test tools in COSMOS to receive bundles, the CLA out port(s) need to be listed in the ports list under the `openc3-cosmos-script-runner-api` section. For example, if you are setting the CLA Out port for contact 0 to 4551 (as in the default configurations), your `compose.yaml` changes should look like this:

```
openc3-cosmos-script-runner-api:
  ports:
    - "4551:4551/udp"
```

To build the `openc3-dtnfsw-plugin`, run the following command from within the `bpnode/cosmos` directory:

```
<Path to COSMOS installation>/openc3.sh cli rake build VERSION=7.0.0
```

COSMOS has a comprehensive [user's guide](#) with a dedicated section for [cFS integration](#), maintained by the OpenC3 team, so further information about the installation and operation of COSMOS is outside the scope of this document.




Figure 11. COSMOS Interface

Appendix B: Host Assumptions

BPNode operates under the following host assumptions:

ADU Proxy

DTN.A.00080

Host provides a mechanism for the DTN Node to exchange ADUs with local client applications.

DTN.A.00090

The mechanism by which ADUs are exchanged between the DTN Node and its client applications is reliable.

Directive Proxy

DTN.A.02120

Host provides a command ingest application that accepts and sends CCSDS telecommand packets in order.

DTN.A.02220

The host will authenticate the source of all directives.

Event Management

DTN.A.00020

Host provides event services.

DTN.A.00030

Host event services provide the capability to filter events by event ID.

DTN.A.00040

Host event services provide the capability to filter events by event type: critical, error, info, and debug.

DTN.A.00050

Host event services include in all event messages the time at which the event occurred.

DTN.A.00060

Host event services will send event messages in a form specified by the flight software framework.

Event Proxy

DTN.A.00020

Host provides event services.

DTN.A.00030

Host event services provide the capability to filter events by event ID.

DTN.A.00040

Host event services provide the capability to filter events by event type: critical, error, info, and

debug.

DTN.A.00050

Host event services include in all event messages the time at which the event occurred.

DTN.A.00060

Host event services will send event messages in a form specified by the flight software framework.

TABLE

DTN.A.02030

Host Table Services supports initialization of the DTN Node configuration.

DTN.A.02031

Host Table Services performs table updates synchronously with the DTN Node application that owns the table to ensure data integrity and table activation atomicity.

DTN.A.02032

Host Table Services supports common viewing of DTN Node application configuration and policy data structures.

DTN.A.02200

The host will load configuration updates via table.

DTN.A.02210

The host will dump DTN Node configuration from tables.

TIME and Time Management

DTN.A.00000

Host provides a monotonic clock that gives a count of ticks since the most recent power-on reset.

DTN.A.00005

Host monotonic time will not rollover during the duration of the mission.

DTN.A.00010

Host monotonic clock reliably counts forward except when reset to zero.

DTN.A.00012

Host provides "host time", an absolute time kept synchronized by the host. Host time is a counter from set epoch known by the DTN Node.

DTN.A.00015

Host determines if host time is valid and provides a time quality indicator.

DTN.A.00017

Host will interface to Position, Navigation, and Timing (PNT) services, if necessary, to synchronize clock time.

Telemetry Proxy

DTN.A.02110

Host provides a telemetry output application for non-DTN communication channels, which transports CCSDS telemetry packets if applicable.

DTN.A.02195

Host provides a method to trigger periodic generation of telemetry.

Appendix C: Acronyms

Acronym

Definition

AA

Application Agent

ADU

Application Data Units

BP

Bundle Protocol

BPA

Bundle Protocol Agent

BPLib

Bundle Protocol Library

BPNode

Bundle Protocol Node

BSC

Bundle Sequence Collection

BSR

Bundle Sequence Range

CBOR

Concise Binary Object Representation

CCS

Compressed Custody Signal

CCSDS

Consultative Committee for Space Data Systems

CT

Custody Transfer

CTDB

Custody Transfer Database

CTEB

Custody Transfer Extension Block

cFE

core Flight Executive

CFDP

CCSDS File Delivery Protocol

cFS

code Flight Systems

CI

Core Infrastructure (in the context of BPNode), Command Ingest (in cFS context)

CL

Convergence Layer

CLA

Convergence Layer Adaptor

DTN

Delay/Disruption Tolerant Networking

EVS

Event Services

FSW

Flight Software

GSFC

Goddard Space Flight Center

OSAL

Operating System Abstraction Layer

PSP

Platform Support Package

RFC

Request for Comments

SB

Software Bus

SC

Stored Commands

TO

Telemetry Output